

Assignment-7: Raster in DBMS, GEO-1006

Ming-Chieh Hu, Carmem E. F. Aires, Benthe Kock

January 2025

Project Description

This report was made by:

1. **Ming-Chieh Hu**, student number: 6186416
2. **Carmem E. F. Aires**, student number: 4325893
3. **Benthe Kock**, student number: 6104819

Contents

1	Check if raster support is available in the database	2
1a	Test extension functionality	2
1b	Raster functions	2
1b.1	ST_Width() and ST_NumBands()	2
1b.2	Other major function categories	2
2	Load raster data in the database	3
2a	Use gdalinfo to retrieve the metadata	3
2b	Completeness of CRS information	4
2c	Load the images in the PostGIS database	5
2d	Comparing the file size	5
2e	Why is it advisable to specify tiling, and how is this related to indexing?	6
2f	Why or why not it would be useful to include one or more ‘overviews’ (-l option) of the raster.	6
3	Retrieve information from, query and manipulate raster data in the database	7
3a	7
3b	7
3c	7
3d	8
3e	8
3f	10
4	Visualize raster data from the database in QGIS frontend	12

1 Check if raster support is available in the database

1a Test extension functionality

We cannot access raster functions directly.

According to the post, postGIS excluded the raster types and functions from the main extension as of version 3.x; a separate CREATE EXTENSION `postgis_raster` is then necessary to get raster support.

```
1 -- Use this query to make the extension available:  
2 create extension postgis_raster;
```

```
1 psql -U username -d dbname -f ./test_raster.sql
```

```
>[~/Documents/TUD/GEO-1006/Assignments/07]$ psql -d geo-1006 -f ./test_raster.sql  
postgis_raster_lib_build_date | postgis_raster_lib_version  
-----+-----  
| 3.5.0 0  
(1 row)  
  
CREATE TABLE  
psql:test_raster.sql:40: NOTICE: SRID value -1 converted to the officially unknown SRID value 0  
INSERT 0 2  
rid  
-----  
1  
2  
(2 rows)  
  
st_width | st_height | st_srid | st_numbands  
-----+-----+-----+-----  
10 | 20 | 0 | 0  
5 | 5 | 0 | 3  
(2 rows)  
  
DROP TABLE
```

1b Raster functions

The ‘test_raster.sql’ script uses some raster functions, like `ST_Width()` and `ST_NumBands()` to test if the raster module works correctly. What type of raster functions are these, and what are the other, major, function categories for raster available in PostGIS?

1b.1 ST_Width() and ST_NumBands()

These 2 functions are both type of **Raster Accessors** according to *postgis.net*. They work the same way as “getter” in other programming language, that is, to get attributes (metadata) from a raster object. Below are their function descriptions:

- `ST_Width()`: Returns the width of the raster in pixels.
- `ST_NumBands()`: Returns the number of bands in the raster object.

1b.2 Other major function categories

The full list of the categories according to *postgis.net* are:

- Raster Management
- Raster Constructors
- Raster Accessors
- Raster Band Accessors
- Raster Pixel Accessors and Setters
- Raster Editors
- Raster Band Editors
- Raster Band Statistics and Analytics
- Raster Inputs

- Raster Outputs
- Raster Processing: Map Algebra
- Built-in Map Algebra Callback Functions
- Raster Processing: DEM (Elevation)
- Raster Processing: Raster to Geometry
- Raster Operators
- Raster and Raster Band Spatial Relationships

In my personal perspective, they can also be divided into 6 major categories:

- Data Setters (constructor)
- Data Getters (accessor)
- Data Conversion Functions (converts object types)
- Data Processing Functions (calculations and processing)
- Data Comparison Functions (Raster Operators)
- Data I/O Functions (save file as a specific format)

2 Load raster data in the database

2a Use gdalinfo to retrieve the metadata

Download the sample images ‘93000_436000.tif’ and ‘93000_436500.tif’ from Brightspace to your local disk. The images are in the GeoTIFF format (Tagged Image File Format) with in some of the tags geo-reference information included. What is the spatial metadata of these images? Use ‘gdalinfo’ to determine this.

93000_436000.tif

```
1 gdalinfo ./93000_436000.tif
```

```
~/Documents/TUD/GEO-1006/Assignments/07]$ gdalinfo ./93000_436000.tif
Driver: GTiff/GeoTIFF
Files: ./93000_436000.tif
Size is 10000, 5000
Coordinate System is:
ENGCRS["unnamed",
    EDATUM["Unknown engineering datum"],
    CS[Cartesian,2],
        AXIS["(E)",east,
            ORDER[1],
            LENGTHUNIT["unknown",1]],
        AXIS["(N)",north,
            ORDER[2],
            LENGTHUNIT["unknown",1]]]
Data axis to CRS axis mapping: 1,2
Origin = (92999.9999999796273,436500.00000000582077)
Pixel Size = (0.100000000000000,-0.100000000000000)
Metadata:
    TIFFTAG_SOFTWARE=Inpho GmbH
    TIFFTAG_XRESOLUTION=1
    TIFFTAG_YRESOLUTION=1
    TIFFTAG_RESOLUTIONUNIT=1 (unitless)
    AREA_OR_POINT=Area
Image Structure Metadata:
    INTERLEAVE=PIXEL
Corner Coordinates:
Upper Left  ( 93000.000, 436500.000)
Lower Left  ( 93000.000, 436000.000)
Upper Right ( 94000.000, 436500.000)
Lower Right ( 94000.000, 436000.000)
Center      ( 93500.000, 436250.000)
Band 1 Block=512x512 Type=Byte, ColorInterp=Red
Band 2 Block=512x512 Type=Byte, ColorInterp=Green
Band 3 Block=512x512 Type=Byte, ColorInterp=Blue
```

93000_436500.tif

```
1 gdalinfo ./93000_436500.tif
```

```
[~/Documents/TUD/GEO-1006/Assignments/07]$ gdalinfo ./93000_436500.tif
Driver: GTiff/GeoTIFF
Files: ./93000_436500.tif
Size is 10000, 5000
Coordinate System is:
ENGCRS["unnamed",
    EDATUM["Unknown engineering datum"],
    CS[Cartesian,2],
        AXIS["(E)",east,
            ORDER[1],
            LENGTHUNIT["unknown",1]],
        AXIS["(N)",north,
            ORDER[2],
            LENGTHUNIT["unknown",1]])
Data axis to CRS axis mapping: 1,2
Origin = (92999.99999999796273,437000.00000000582077)
Pixel Size = (0.100000000000000,-0.100000000000000)
Metadata:
    TIFFTAG_SOFTWARE=Inpho GmbH
    TIFFTAG_XRESOLUTION=1
    TIFFTAG_YRESOLUTION=1
    TIFFTAG_RESOLUTIONUNIT=1 (unitless)
    AREA_OR_POINT=Area
Image Structure Metadata:
    INTERLEAVE=PIXEL
Corner Coordinates:
Upper Left  (  93000.000,  437000.000)
Lower Left  (  93000.000,  436500.000)
Upper Right (  94000.000,  437000.000)
Lower Right (  94000.000,  436500.000)
Center      (  93500.000,  436750.000)
Band 1 Block=512x512 Type=Byte, ColorInterp=Red
Band 2 Block=512x512 Type=Byte, ColorInterp=Green
Band 3 Block=512x512 Type=Byte, ColorInterp=Blue
```

According to the displayed information, we're able to get a grasp on these 2 images:

- The size for both the images are 10000 by 5000 (width, height).
- They both have 3 bands (RGB)
- If we tile both images, 93000_436500.tif will be in the north of 93000_436000.tif

2b Completeness of CRS information

Is the spatial reference information of the images complete? (the coordinates in the GeoTIFF files refer to the Dutch national grid, Spatial Reference ID = 28992)

CRS Information from gdalinfo

```
1 Coordinate System is:
2 ENGCRS["unnamed",
3     EDATUM["Unknown engineering datum"],
4     CS[Cartesian,2],
5         AXIS["(E)",east,
6             ORDER[1],
7             LENGTHUNIT["unknown",1]],
8         AXIS["(N)",north,
9             ORDER[2],
10            LENGTHUNIT["unknown",1]])
11 Data axis to CRS axis mapping: 1,2
12 Origin = (92999.99999999796273,437000.00000000582077)
```

The CRS metedata for these 2 images are not very clear, specifically:

Coordinate System

- ENGCRS ["unnamed"] : Indicates that the dataset uses an undefined CRS or non-georeferenced system.
- EDATUM ["Unknown engineering datum"] : The datum is not defined.

- CS[Cartesian,2]: Specifies a 2D Cartesian coordinate system with two axes (East and North).

Axes

- AXIS["(E)",east]: First axis is labeled East.
- AXIS["(N)",north]: Second axis is labeled North.
- Both axes have LENGTHUNIT["unknown",1], meaning the unit of measurement for the coordinates is undefined. It could be meters, feet, or any other unit.

Data Axis to CRS Axis Mapping

- 1,2: Confirms that the dataset's data axis aligns directly with the CRS axis (East → 1, North → 2).

2c Load the images in the PostGIS database

Using `raster2pgsql` to load raster data involves a 2-step procedure. In the first step `raster2pgsql` analyzes the input images and processes the options specified. This results in a file with SQL commands which have to be executed in a second step to actually load the raster data in the database. For this to work properly you have to run the `raster2pgsql` command from the operating system “command line”.

Step 1

```
1 raster2pgsql -C -s 28992 -t 500x500 -I -Y ./*.tif rotterdam > ./load2.sql
```

Step 2

```
1 psql -U username -d dbname -f ./load2.sql
```

```
[~/Documents/TUD/GEO-1006/Assignments/07]$ psql -U mchu -d geo-1006 -f ./load2.sql
BEGIN
CREATE TABLE
COPY 50
CREATE INDEX
ANALYZE
psql:load2.sql:421: NOTICE: Adding SRID constraint
psql:load2.sql:421: NOTICE: Adding scale-X constraint
psql:load2.sql:421: NOTICE: Adding scale-Y constraint
psql:load2.sql:421: NOTICE: Adding blocksize-X constraint
psql:load2.sql:421: NOTICE: Adding blocksize-Y constraint
psql:load2.sql:421: NOTICE: Adding alignment constraint
psql:load2.sql:421: NOTICE: Adding number of bands constraint
psql:load2.sql:421: NOTICE: Adding pixel type constraint
psql:load2.sql:421: NOTICE: Adding nodata value constraint
psql:load2.sql:421: NOTICE: Adding out-of-database constraint
psql:load2.sql:421: NOTICE: Adding maximum extent constraint
    addresserconstraints
-----
t
(1 row)
COMMIT
```

2d Comparing the file size

After loading the images check and report on the size of the data: how does the size of the ‘raw’ images compare to the size of the same data stored in the database? Can you explain the differences, if there are any?

File on the disk

```
1 ls -l
```

```
[~/Documents/TUD/GEO-1006/Assignments/07]$ ls -l
total 1845552
-rw-r--r--@ 1 mchu  staff  157288404 Jan  7 12:54 93000_436000.tif
-rw-r--r--@ 1 mchu  staff  157288404 Jan  7 12:54 93000_436500.tif
```

The file on the disk is approximately 314 MB (157 megabytes each).

File in the DB

```
1 select pg_size.pretty(pg_table_size('rotterdam'));
```

	ABC pg_size.pretty
1	217 MB

The file in the database is approximately 217 MB, with 400 rows (because we use tile size of 500x500) in total.

Comparing differences

The main differences are in:

- Tiling Improves Compression Efficiency:
 - Splitting the raster into 500x500 tiles enhances compression by exploiting repetitive patterns within smaller, localized areas.
 - Uniform terrain areas particularly benefit from this approach.
- Efficient Storage of Numeric Data:
 - PostgreSQL uses a structured and compact format for raster storage.
 - Redundant parts of the TIFF file structure (e.g. headers) are optimized or removed during import.

2e Why is it advisable to specify tiling, and how is this related to indexing?

Tiling divides the data into chunks, which is good for data performance. This approach is particularly useful for large datasets. When loading the data, only the visible tiles are retrieved. Similarly, when querying, only the number of tiles required will be retrieved, depending on the region required. Indexing is required to retrieve the data. Each tile, when created, manages a number of x and y pixels that are stored in each BLOB field.

2f Why or why not it would be useful to include one or more ‘overviews’ (-l option) of the raster.

An overview is a low-resolution representation of the core tables of a raster. This allows quick map zoom-in and-out. Computations can also be performed on them in less time. However, their use depends on the application, as each pixel generalises a larger area, they will give less accurate results than the same calculation performed in their higher resolution parents.

3 Retrieve information from, query and manipulate raster data in the database

3a

```
1 select ST_Width(rast), ST_Height(rast), ST_SRID(rast), ST_NumBands(rast), ST_GeoReference(rast)
2 from rast_table rt;
```

	123 st_width	123 st_height	123 st_srid	123 st_numbands	st_georeference
1	500	500	28,992	3	0.1000000000 0.0000000000 -0.1000000000 92999.999999998 436500.0000000010
2	500	500	28,992	3	0.1000000000 0.0000000000 -0.1000000000 93049.999999998 436500.0000000010
3	500	500	28,992	3	0.1000000000 0.0000000000 -0.1000000000 93099.999999998 436500.0000000010
4	500	500	28,992	3	0.1000000000 0.0000000000 -0.1000000000 93149.999999998 436500.0000000010
5	500	500	28,992	3	0.1000000000 0.0000000000 -0.1000000000 93199.999999998 436500.0000000010
6	500	500	28,992	3	0.1000000000 0.0000000000 -0.1000000000 93249.999999998 436500.0000000010
7	500	500	28,992	3	0.1000000000 0.0000000000 -0.1000000000 93299.999999998 436500.0000000010
8	500	500	28,992	3	0.1000000000 0.0000000000 -0.1000000000 93349.999999998 436500.0000000010
9	500	500	28,992	3	0.1000000000 0.0000000000 -0.1000000000 93399.999999998 436500.0000000010
10	500	500	28,992	3	0.1000000000 0.0000000000 -0.1000000000 93349.999999998 436500.0000000010
11	500	500	28,992	3	0.1000000000 0.0000000000 -0.1000000000 93499.999999998 436500.0000000010
12	500	500	28,992	3	0.1000000000 0.0000000000 -0.1000000000 93549.999999998 436500.0000000010
13	500	500	28,992	3	0.1000000000 0.0000000000 -0.1000000000 93599.999999998 436500.0000000010
14	500	500	28,992	3	0.1000000000 0.0000000000 -0.1000000000 93649.999999998 436500.0000000010
15	500	500	28,992	3	0.1000000000 0.0000000000 -0.1000000000 93699.999999998 436500.0000000010
16	500	500	28,992	3	0.1000000000 0.0000000000 -0.1000000000 93749.999999998 436500.0000000010
17	500	500	28,992	3	0.1000000000 0.0000000000 -0.1000000000 93799.999999998 436500.0000000010
18	500	500	28,992	3	0.1000000000 0.0000000000 -0.1000000000 93849.999999998 436500.0000000010
19	500	500	28,992	3	0.1000000000 0.0000000000 -0.1000000000 93899.999999998 436500.0000000010
20	500	500	28,992	3	0.1000000000 0.0000000000 -0.1000000000 93949.999999998 436500.0000000010

3b

```
1 select rid, ST_Value(rast, 1, ST_SetSRID(ST_Point(93780, 436208), 28992)) as band1,
2   → ST_Value(rast, 2, ST_SetSRID(ST_Point(93780, 436208), 28992)) as band2, ST_Value(rast, 3,
3   → ST_SetSRID(ST_Point(93780, 436208), 28992))as band3
4 from rast_table rt
5 where ST_Intersects(rast, ST_SetSRID(ST_Point(93780, 436208), 28992));
```

	123 rid	123 band1	123 band2	123 band3
1	116	130	124	105

3c

```
1 -- This doesn't save the value
2 select ST_SetValue(rast, 1, ST_SetSRID(ST_Point(93780, 436208), 28992), 101)
3 from rast_table rt
4 where ST_Intersects(rast, ST_SetSRID(ST_Point(93780, 436208), 28992));
5
6 -- This one will save the value
7 UPDATE rast_table
8 SET rast = ST_SetValue(
9   ST_SetValue(
10     ST_SetValue(rast, 1, ST_SetSRID(ST_Point(93780, 436208), 28992), 101),-- Band 1
11     2, ST_SetSRID(ST_Point(93780, 436208), 28992), 102),-- Band 2
12     3, ST_SetSRID(ST_Point(93780, 436208), 28992), 103)-- Band 3
13 WHERE ST_Intersects(rast, ST_SetSRID(ST_Point(93780, 436208), 28992));
```

	123 rid	123 band1	123 band2	123 band3
1	116	101	102	103

3d

```
1 CREATE TABLE clipped_raster_3d AS
2 select ST_Clip(rast, ST_MakeEnvelope(93750,436186, 93778,436211, 28992))
3 from rast_table rt
4 where ST_Intersects(rast, ST_MakeEnvelope(93750,436186, 93778,436211, 28992));
```



Figure 14: 3d - Clipped raster

3e

```
1 CREATE TABLE clipped_raster_3e AS
2 SELECT ST_Clip(rast, ST_MakeEnvelope(93550,436350, 93750,436600, 28992)) AS rast
3 FROM rotterdam rt
4 WHERE ST_Intersects(rast, ST_MakeEnvelope(93550,436350, 93750,436600, 28992));
```



Figure 16: 3e - extent view in QGIS



Figure 17: 3e - Clipped raster

3f

```
1 CREATE TABLE rescaled_raster AS
2 SELECT
3     ST_Rescale(
4         ST_Clip(rast, ST_MakeEnvelope(93550, 436350, 93750, 436600, 28992)),
5         0.25,
6         0.25
7     ) AS rast
8 FROM rotterdam rt
9 WHERE ST_Intersects(rast, ST_MakeEnvelope(93550, 436350, 93750, 436600, 28992));
```



Figure 19: 3f - Rescaled raster

General

Name clipped_raster_3e
Source dbname='Assignment7' host=localhost port=5432 srid=0 table="public"."clipped_raster_3e" (rast)
Provider postgresraster

Information from provider

Extent 93549.9999999997962732,436349.9000000006053597 : 93750.0999999998020940,436600.0000000005820766
Width 2001
Height 2501
Data type Byte - Eight bit unsigned integer
Additional information
Is Tiled false
Overviews
Pixel Size 0.1, -0.1
Primary Keys SQL ctid
Temporal Column
Where Clause SQL

Figure 20: 3f - Properties before rescale

General

Name rescaled_raster
Source dbname='Assignment7' host=localhost port=5432 srid=0 table="public"."rescaled_raster" (rast)
Provider postgresraster

Information from provider

Extent 93549.9999999997962732,436349.7500000005820766 : 93750.2499999997962732,436600.0000000005820766
Width 801
Height 1001
Data type Byte - Eight bit unsigned integer
Additional information
Is Tiled false
Overviews
Pixel Size 0.25, 0.25
Primary Keys SQL ctid
Temporal Column
Where Clause SQL

Figure 21: 3f - Properties after rescale

4 Visualize raster data from the database in QGIS frontend

We have presented some of the results of Part III in QGIS and attached them to Part 3.

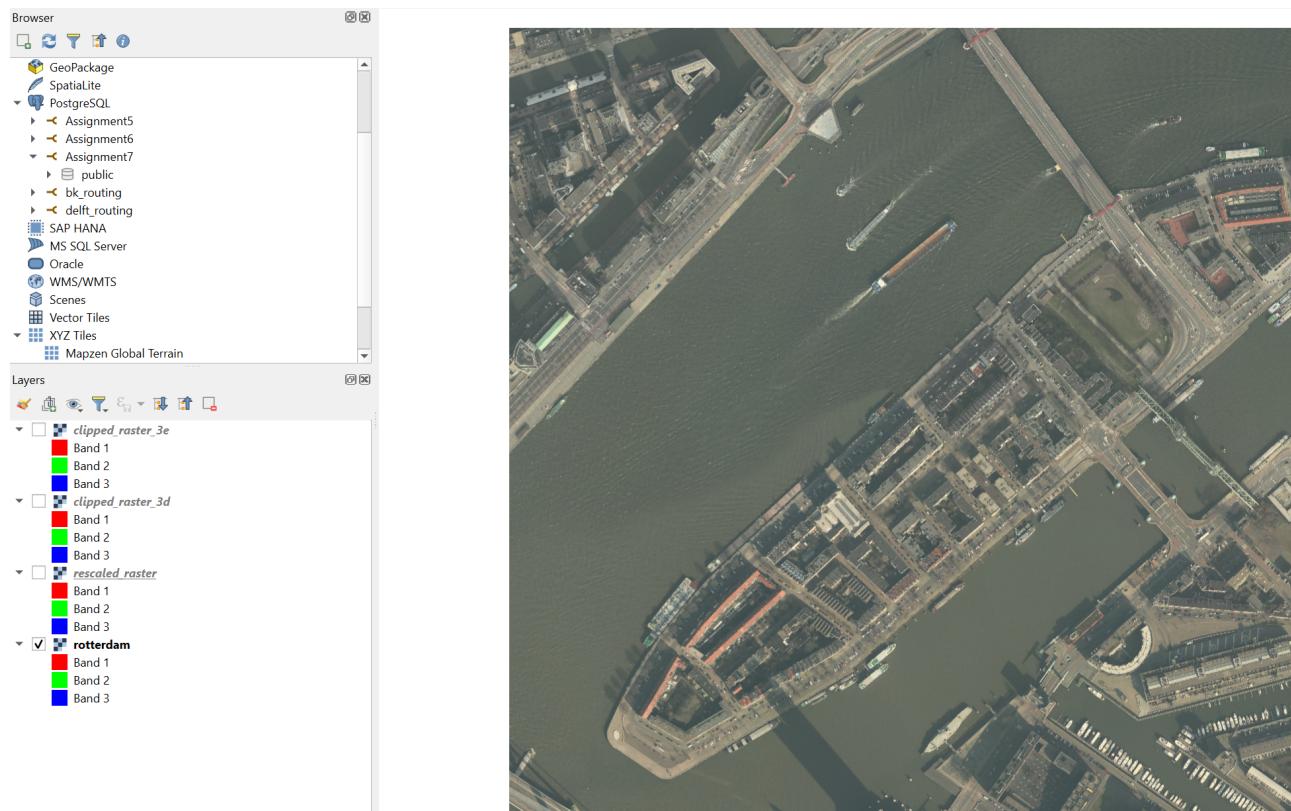


Figure 22: Overview of the raster data in QGIS